

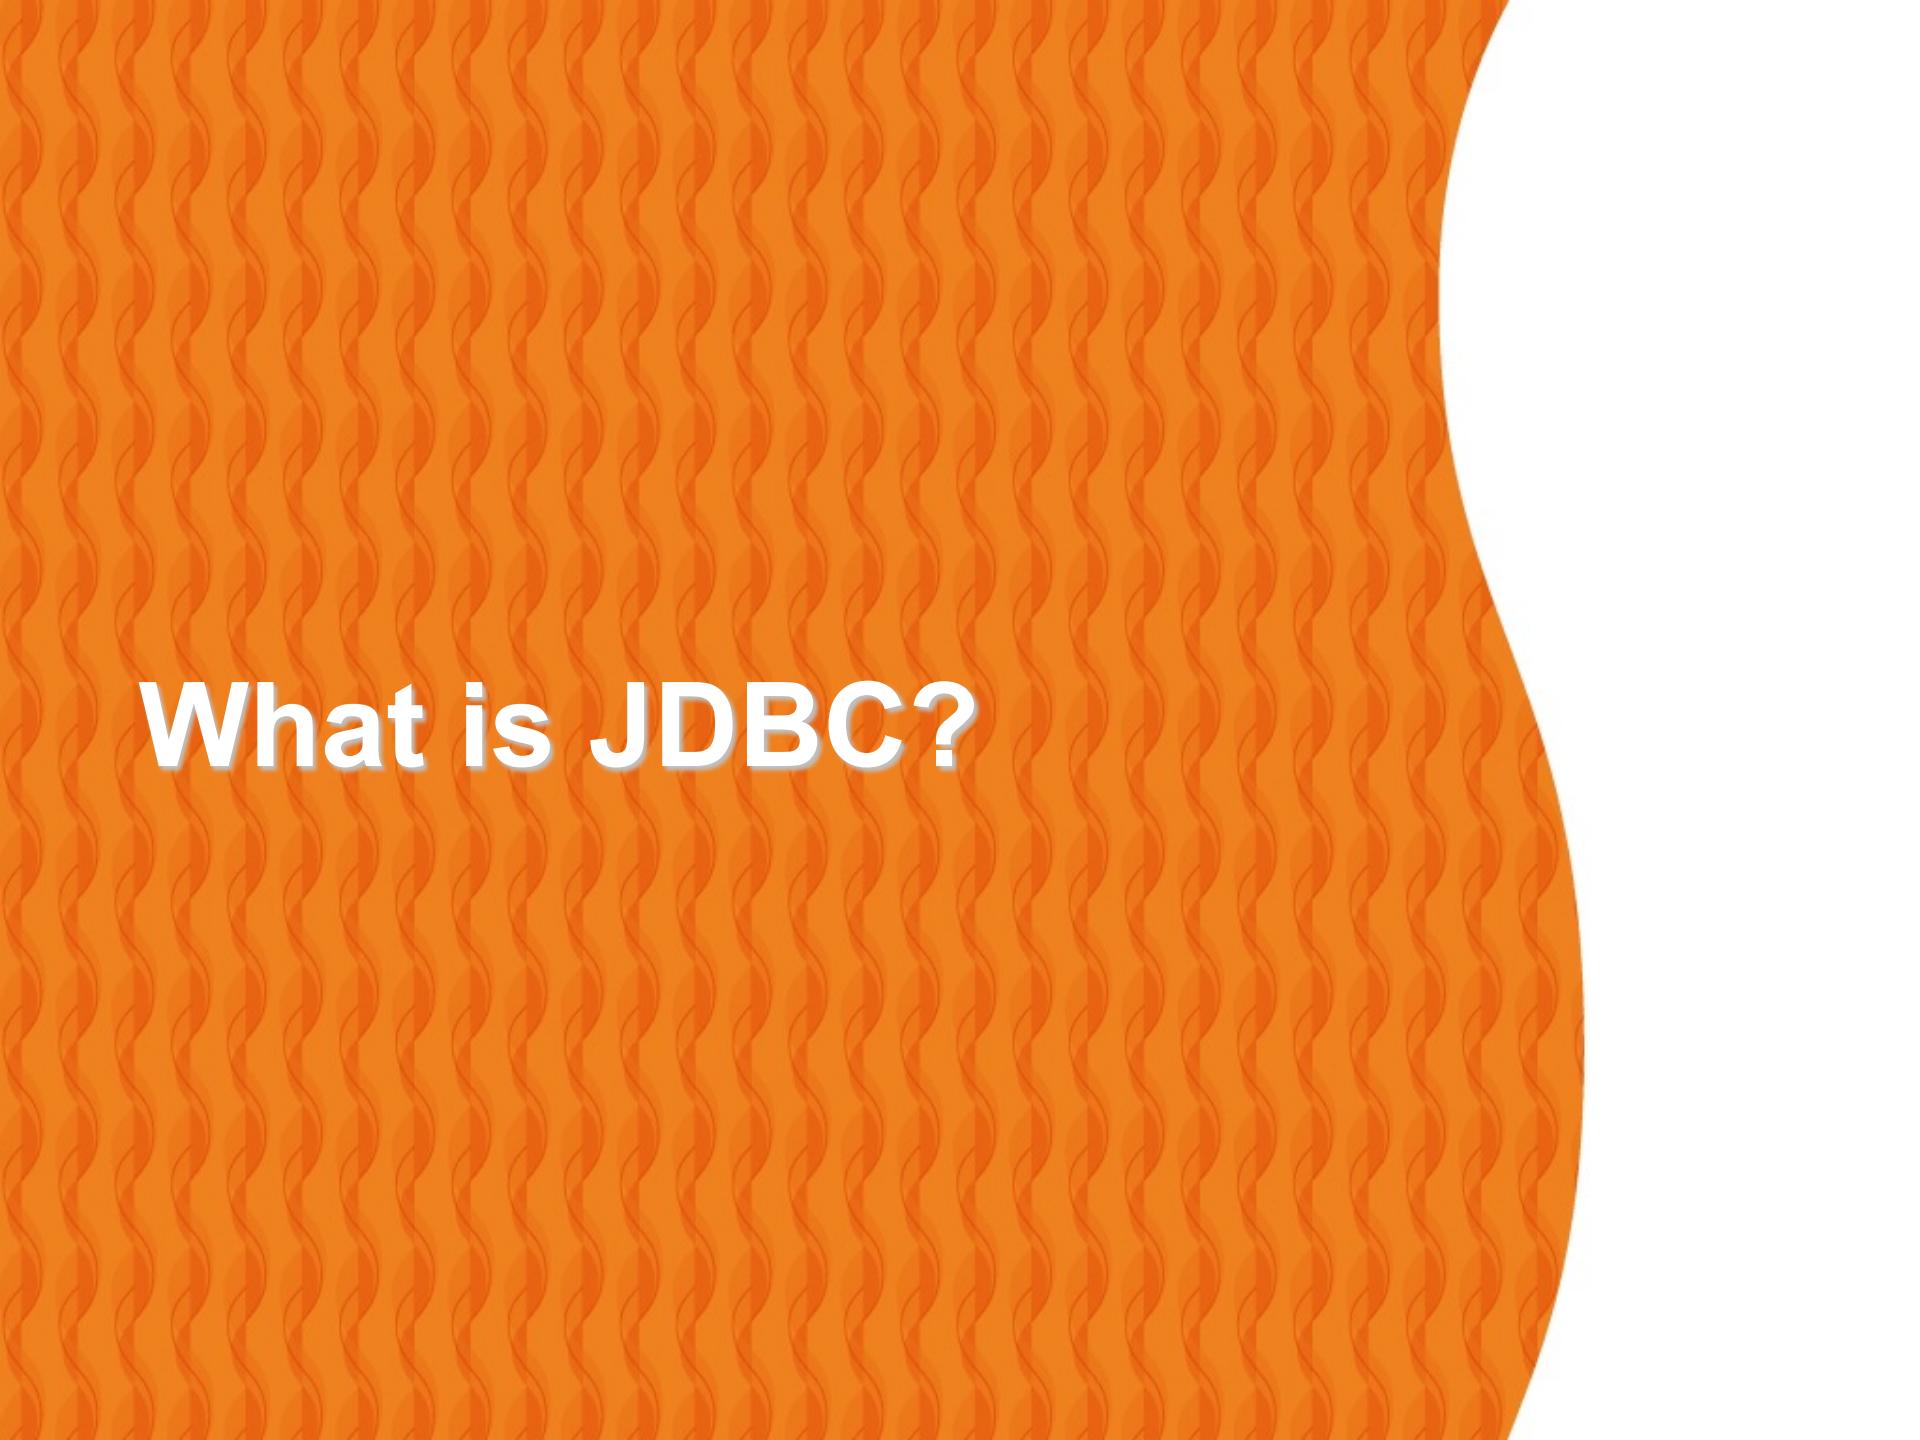
# JDBC Overview

**“Code with Passion!”**



# Topics

- What is JDBC?
- Database URL
- Step By Step Usage of JDBC API
- PreparedStatement
- CallableStatement
- Transaction
- Transaction Isolation levels
- DataSource



# What is JDBC?

# What is JDBC?

- ◆ Standard Java API for accessing relational database
  - Hides database specific details from application
- ◆ Part of Java SE (J2SE)
  - Java SE 6+ has JDBC 4

# JDBC API

- Defines a set of Java Interfaces, which are implemented by vendor-specific JDBC client drivers
  - > Applications use this set of Java interfaces for performing database operations – portability is guaranteed
- Majority of JDBC API is located in `java.sql` package
  - DriverManager, Connection, ResultSet, DatabaseMetaData, ResultSetMetaData, PreparedStatement, CallableStatement and Types
- Other advanced functionality exists in the `javax.sql` package
  - DataSource

# What is JDBC Driver?

- Database specific implementation of JDBC interfaces
  - > Each database has its own JDBC driver implementing the same JDBC interfaces
- Example JDBC drivers
  - > mysql-connector-java-5.1.34.jar (for MySQL)
  - > derbyclient-10.8.1.2.jar (for JavaDB)

# Database URL

# Database URL

- Used to make a connection to the database
- Can contain server, port, protocol etc
- Format
  - > `jdbc:[subprotocol]:[server-location]/[database-name]`
- Examples
  - > `jdbc:mysql://localhost:3306/mydatabase`
  - > `jdbc:derby://localhost:1527/sample`

# **Step by Step Usage of JDBC API**

# Steps of Using JDBC

1. Load database-specific JDBC driver
2. Get a Connection object
3. Get a Statement object
4. Execute queries and/or updates
5. Read results
6. Read Meta-data (optional step)
7. Close Statement and Connection objects

# 1. Load DB-Specific Database Driver

- To manually load the database driver and register it with the [DriverManager](#), load its class file:  
Class.forName(<database-driver>)

```
try {  
    // This loads an instance of the MySQL Driver.  
    // The driver has to be in the classpath.  
    Class.forName("com.mysql.jdbc.Driver");  
} catch (ClassNotFoundException cnfe){  
    System.out.println("'" + cnfe);  
}
```

## 2. Get a Connection Object

- **DriverManager** class is responsible for selecting the database and creating the database connection
- Create the database connection as follows:

```
try {  
    Connection connection =  
        DriverManager.getConnection("jdbc:mysql://localhost:3306/mydb",  
            "username", "password");  
}  
} catch(SQLException sqle) {  
    System.out.println("") + sqle);  
}
```

# DriverManager & Connection classes

- java.sql.DriverManager
  - > getConnection(String url, String user, String password) throws SQLException
- java.sql.Connection
  - > Statement createStatement() throws SQLException
  - > void close() throws SQLException
  - > void setAutoCommit(boolean b) throws SQLException
  - > void commit() throws SQLException
  - > void rollback() throws SQLException

### 3. Get a Statement Object

- Create a Statement Object from Connection object
  - > *Statement statement = connection.createStatement();*
- The Statement object is used for executing a query statement

## 4. Executing Query or Update

- Using the Statement object, you can do  
SELECT/UPDATE/DELETE
- SELECT
  - > *ResultSet rs = statement.executeQuery("select \* from customer\_tbl");*
- UPDATE
  - > *int iReturnValue = statement.executeUpdate("update manufacture\_tbl set name = ABC where mfr\_num = 19985678");*

## 5. Reading Results

- Once you have *ResultSet*, you can easily access the data by looping through it
- The iterator is initialized to a position before the first row so you must call *next()* once to move it to the first row

```
while (rs.next()){  
    // Wrong this will generate an error since column index starts from 1  
    String value0 = rs.getString(0);  
  
    // Correct!  
    String value1 = rs.getString(1);  
    int   value2 = rs.getInt(2);  
    int   value3 = rs.getInt("ADDR_LN1");  
}
```

## 5. Reading Results (Continued)

- When retrieving data from the *ResultSet*, use the appropriate `getXXX()` method
- There is an appropriate `getXXX()` method of each `java.sql.Types` datatype
  - > `getString()`
  - > `getInt()`
  - > `getDouble()`
  - > `getObject()`
  - > ...

## 6.a Read ResultSet MetaData (Optional)

- Once you have the *ResultSet*, you can obtain the meta data about the query
  - > *ResultSetMetaData rsMeta = rs.getMetaData();*
- From the *ResultSetMetaData*, you get meta data such as
  - > Column count
  - > Column name and type

# ResultSetMetaData Example

```
ResultSetMetaData meta = rs.getMetaData();
// Get the column count
int iColumnCount = meta.getColumnCount();

for (int i =1 ; i <= iColumnCount ; i++){
    System.out.println("Column Name: " + meta.getColumnName(i));
    System.out.println("Column Type" + meta.getColumnType(i));
    System.out.println("Display Size: " + meta getColumnDisplaySize(i) );
    System.out.println("Precision: " + meta.getPrecision(i));
    System.out.println("Scale: " + meta.getScale(i) );
}
```

## 6.b Read DataBaseMetaData (Optional)

- Once you have the *Connection* object, you can obtain the Meta Data about the database
  - > *DatabaseMetaData dbmetadata = connection.getMetaData();*
- There are approximately 150 methods in the *DatabaseMetaData* class.

# PreparedStatement

# What is PreparedStatement?

- Represents a precompiled SQL statement.
- A SQL statement is precompiled and stored in a *PreparedStatement* object
- This object can then be used to efficiently execute this statement multiple times.

# Example: PreparedStatement

```
PreparedStatement pstmt =
```

```
    con.prepareStatement("UPDATE EMPLOYEES  
        SET SALARY = ? WHERE ID = ?");
```

```
pstmt.setBigDecimal(1, 100000.00)
```

```
pstmt.setInt(2, 110592)
```

```
pstmt.executeUpdate();
```

```
pstmt.setBigDecimal(1, 200000.00)
```

```
pstmt.setInt(2, 222222)
```

```
pstmt.executeUpdate();
```

# **CallableStatement**

# CallableStatement

- Used to execute SQL stored procedures

# CallableStatement cont.

- A *CallableStatement* object contains a call to a stored procedure; it does not contain the stored procedure itself.
- The first line of code below creates a call to the stored procedure SHOW\_SUPPLIERS using the connection con .
- The part that is enclosed in curly braces is the escape syntax for stored procedures.

```
CallableStatement cs = con.prepareCall("{call SHOW_SUPPLIERS}");  
ResultSet rs = cs.executeQuery();
```

# CallableStatement Example

Here is an example using IN, OUT and INOUT parameters

```
// set int IN parameter  
cstmt.setInt( 1, 333 );  
// register int OUT parameter  
cstmt.registerOutParameter( 2, Types.INTEGER );  
// set int INOUT parameter  
cstmt.setInt( 3, 666 );  
// register int INOUT parameter  
cstmt.registerOutParameter( 3, Types.INTEGER );  
// You then execute the statement with no return value  
cstmt.execute(); // could use executeUpdate()  
// get int OUT and INOUT  
int iOUT = cstmt.getInt( 2 );  
int iINOUT = cstmt.getInt( 3 );
```

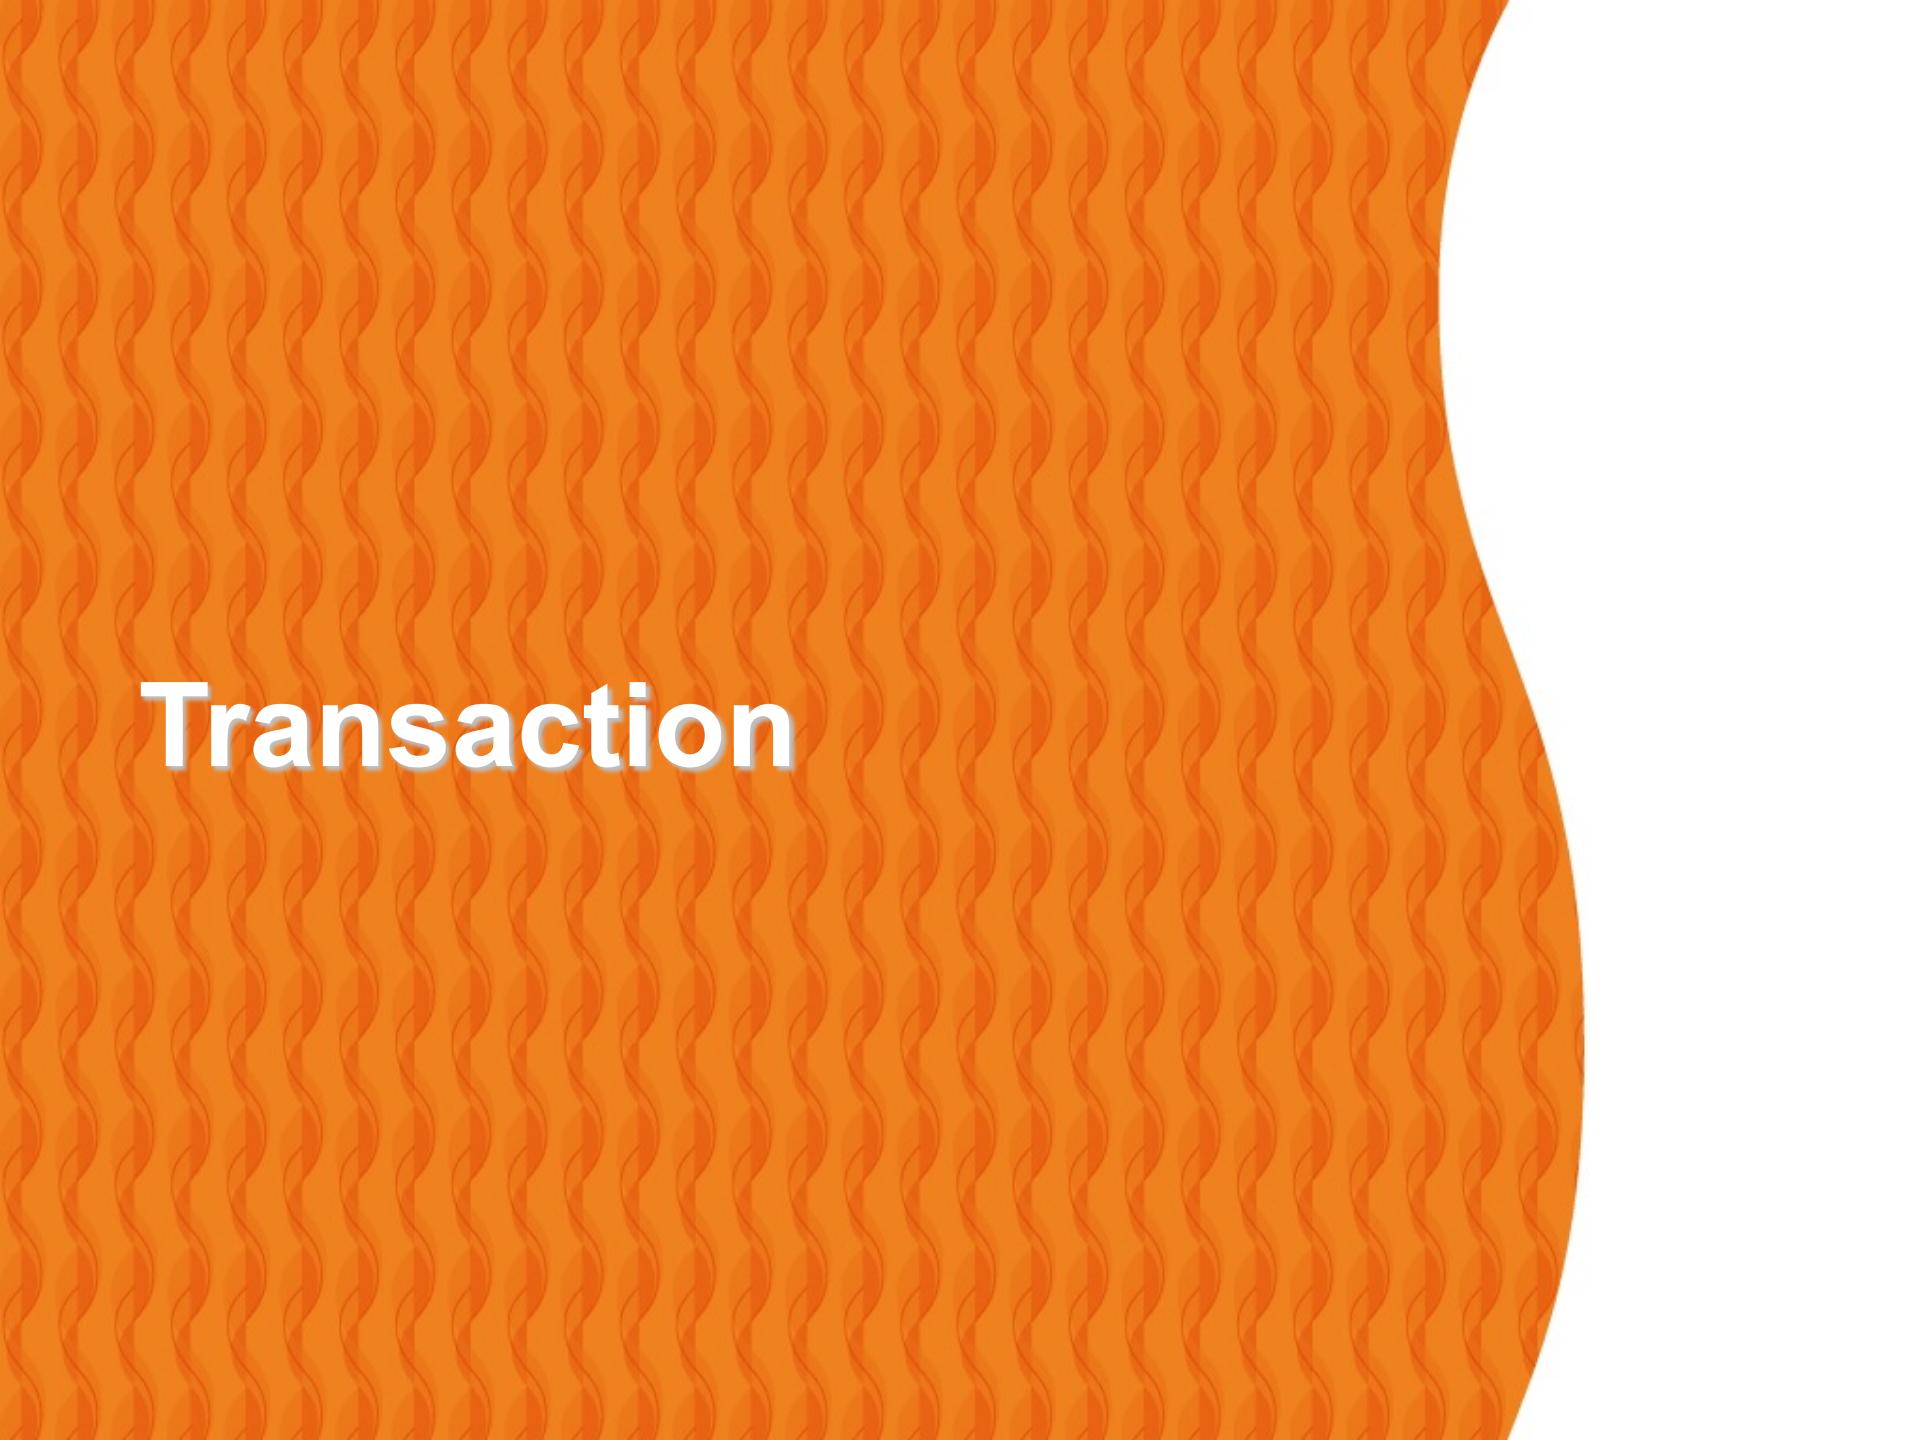
# Stored Procedure example

```
FUNCTION event_list (appl_id_in  VARCHAR2,  
                    dow_in      VARCHAR2,  
                    event_type_in VARCHAR2 OUT,  
                    status_in    VARCHAR2 INOUT)  
  
RETURN ref_cur;
```

# Lab:

**Exercise 1: Single Table JDBC**  
**Exercise 2: Joined Table JDBC**  
**Exercise 3: Advanced JDBC**  
**1020\_jdbc\_overview.zip**





# Transaction

# Why use Transaction?

- Committing of each statement, which is a default behavior, could be very time consuming
  - > This is the default behavior – *AutoCommit* is set to true as a default
- By setting *AutoCommit* to *false*, the developer can update the database more than once without committing and then commit the entire transaction as a whole
- Also, if each statement is dependant on the other, the entire transaction can be rolled back and the user notified

# JDBC Transaction Methods

- `setAutoCommit()`
  - > If set true, every executed statement is committed immediately
- `commit()`
  - > Relevant only if `setAutoCommit(false)`
  - > Commit operations performed since the opening of a Connection or last `commit()` or `rollback()` calls
- `rollback()`
  - > Relevant only if `setAutoCommit(false)`
  - > Cancels all operations performed

# Transactions Example

```
String url = "jdbc:mysql://localhost/testdb";
String username = "root";
String password = "";
Class.forName("com.mysql.jdbc.Driver");
Connection conn = null;
try {
    conn = DriverManager.getConnection(url, username, password);
    conn.setAutoCommit(false);

    Statement st = conn.createStatement();
    st.execute("INSERT INTO orders (username, order_date) VALUES ('java', '2007-12-13')",
              Statement.RETURN_GENERATED_KEYS);

    ResultSet keys = st.getGeneratedKeys();
    int id = 1;
    while (keys.next()) {
        id = keys.getInt(1);
    }

// To be continued
```

# Transactions Example

```
PreparedStatement pst = conn.prepareStatement("INSERT INTO order_details (order_id, product_id,
    quantity, price) VALUES (?, ?, ?, ?)");
pst.setInt(1, id);
pst.setString(2, "1");
pst.setInt(3, 10);
pst.setDouble(4, 100);
pst.execute();

conn.commit();
System.out.println("Transaction commit...");
} catch (SQLException e) {
if (conn != null) {
    conn.rollback();
    System.out.println("Connection rollback... ");
}
e.printStackTrace();
} finally {
if (conn != null && !conn.isClosed()) {
    conn.close();
}
}
```

# Lab:

**Exercise 4: Transaction  
1020\_jdbc\_overview.zip**



# Transaction Isolation Levels

# 4 Isolation Levels

- Level 1 - READ UNCOMMITTED
  - > Provides lowest level of isolation among transactions but best performing
- Level 2 - READ COMMITTED
- Level 3 - REPEATABLE READ
- Level 4 - SERIALIZABLE
  - > Provides the highest level of isolation among transactions but least performing

# Level 1 - READ UNCOMMITTED Isolation level

- Causes 'dirty reads' symptom
  - > Uncommitted changes in one transaction (client #1 below) is visible in other transactions (client #2 below)

Client #1: Start Transaction---INSERT a record A-----ROLLBACK----->

Client #2: Start Transaction-----**(See Record A:Dirty read)**-----**(Does not see Record A)**-->

## Level 2 - READ COMMITTED Isolation level

- Committed updates in one transaction (client # 1 below) are visible within another transaction (client #2 below). This means identical queries within a transaction (in client #2 below) can return differing results

Client #1: Start Transaction---INSERT a record A-----COMMIT----->

Client #2: Start Transaction-----**(Not see Record A)**-----**(See Record A)**-----COMMIT----->

## Level 3 - REPEATABLE READ Isolation level

- Committed changes in one transaction (client #1 below) is visible in another transaction (client #2 below) only after its own Commit.
  - > Within a transaction, all reads are consistent.
- The default isolation level for most databases

Client #1: Start Transaction---INSERT a record A-----COMMIT----->

Client #2: Start Transaction-----(Not see Record A)-----COMMIT--(See Record A) --->

## Level 4 - SERIALIZABLE Isolation level

- Transactions are serialized
- Until the previous transaction ends via either COMMIT or ROLLBACK, the database operations in other transactions are blocked
- Highest isolation level but not practical in real-life environment

Client #1: Start Transaction---INSERT a record A-----COMMIT----->

Client #2: Start Transaction-----**(SELECT Blocked)**----- **(SELECT returns)**--COMMIT-->

# Lab:

**Exercise 7: Check Isolation Level  
1020\_jdbc\_overview.zip**



# Code with Passion!

